

```
In [1]: import pandas as pd
import numpy as np
import warnings
warnings.filterwarnings("ignore")
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.metrics import accuracy_score, classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split
```

```
In [2]: data=pd.read_csv(r"C:\Users\Kishore\OneDrive\Desktop\CSV Files\bankloan.csv")
```

```
In [3]: data
```

```
Out[3]:
```

	ID	Age	Experience	Income	ZIP.Code	Family	CCAvg	Education	Mortgage	Personal.Loan	Se
0	1	25	1	49	91107	4	1.6	1	0	0	
1	2	45	19	34	90089	3	1.5	1	0	0	
2	3	39	15	11	94720	1	1.0	1	0	0	
3	4	35	9	100	94112	1	2.7	2	0	0	
4	5	35	8	45	91330	4	1.0	2	0	0	
...	
4995	4996	29	3	40	92697	1	1.9	3	0	0	
4996	4997	30	4	15	92037	4	0.4	1	85	0	
4997	4998	63	39	24	93023	2	0.3	3	0	0	
4998	4999	65	40	49	90034	3	0.5	2	0	0	
4999	5000	28	4	83	92612	3	0.8	1	0	0	

5000 rows × 14 columns



```
In [4]: data.shape
```

```
Out[4]: (5000, 14)
```

In [5]: data.describe()

Out[5]:

	ID	Age	Experience	Income	ZIP.Code	Family	CCAvg	Ed
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	45.338400	20.104600	73.774200	93152.503000	2.396400	1.937938	1.000000
std	1443.520003	11.463166	11.467954	46.033729	2121.852197	1.147663	1.747659	0.000000
min	1.000000	23.000000	-3.000000	8.000000	9307.000000	1.000000	0.000000	1.000000
25%	1250.750000	35.000000	10.000000	39.000000	91911.000000	1.000000	0.700000	1.000000
50%	2500.500000	45.000000	20.000000	64.000000	93437.000000	2.000000	1.500000	2.000000
75%	3750.250000	55.000000	30.000000	98.000000	94608.000000	3.000000	2.500000	3.000000
max	5000.000000	67.000000	43.000000	224.000000	96651.000000	4.000000	10.000000	3.000000

In [6]: data.head(5)

Out[6]:

	ID	Age	Experience	Income	ZIP.Code	Family	CCAvg	Education	Mortgage	Personal.Loan	Securities
0	1	25	1	49	91107	4	1.6	1	0	0	0
1	2	45	19	34	90089	3	1.5	1	0	0	0
2	3	39	15	11	94720	1	1.0	1	0	0	0
3	4	35	9	100	94112	1	2.7	2	0	0	0
4	5	35	8	45	91330	4	1.0	2	0	0	0

In [7]: data.tail(5)

Out[7]:

	ID	Age	Experience	Income	ZIP.Code	Family	CCAvg	Education	Mortgage	Personal.Loan	Securities
4995	4996	29	3	40	92697	1	1.9	3	0	0	0
4996	4997	30	4	15	92037	4	0.4	1	85	0	0
4997	4998	63	39	24	93023	2	0.3	3	0	0	0
4998	4999	65	40	49	90034	3	0.5	2	0	0	0
4999	5000	28	4	83	92612	3	0.8	1	0	0	0

In [8]: data.shape

Out[8]: (5000, 14)

```
In [9]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 14 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   ID                    5000 non-null   int64  
 1   Age                   5000 non-null   int64  
 2   Experience             5000 non-null   int64  
 3   Income                5000 non-null   int64  
 4   ZIP.Code              5000 non-null   int64  
 5   Family                5000 non-null   int64  
 6   CCAvg                 5000 non-null   float64 
 7   Education             5000 non-null   int64  
 8   Mortgage              5000 non-null   int64  
 9   Personal.Loan         5000 non-null   int64  
10   Securities.Account     5000 non-null   int64  
11   CD.Account            5000 non-null   int64  
12   Online                5000 non-null   int64  
13   CreditCard            5000 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 547.0 KB
```

```
In [10]: data.isnull().sum()
```

```
Out[10]: ID                0
Age                0
Experience         0
Income            0
ZIP.Code          0
Family            0
CCAvg             0
Education         0
Mortgage          0
Personal.Loan     0
Securities.Account 0
CD.Account        0
Online            0
CreditCard       0
dtype: int64
```

```
In [11]: data.min()
```

```
Out[11]: ID                1.0
Age                23.0
Experience        -3.0
Income            8.0
ZIP.Code         9307.0
Family            1.0
CCAvg             0.0
Education         1.0
Mortgage          0.0
Personal.Loan     0.0
Securities.Account 0.0
CD.Account        0.0
Online            0.0
CreditCard       0.0
dtype: float64
```

In [12]: `data.max()`

```
Out[12]: ID                5000.0
Age                67.0
Experience          43.0
Income             224.0
ZIP.Code           96651.0
Family             4.0
CCAvg              10.0
Education           3.0
Mortgage           635.0
Personal.Loan       1.0
Securities.Account  1.0
CD.Account          1.0
Online             1.0
CreditCard         1.0
dtype: float64
```

In [13]: `list(data)`

```
Out[13]: ['ID',
'Age',
'Experience',
'Income',
'ZIP.Code',
'Family',
'CCAvg',
'Education',
'Mortgage',
'Personal.Loan',
'Securities.Account',
'CD.Account',
'Online',
'CreditCard']
```

In [14]: `data1 = data.drop(['Securities.Account', 'ZIP.Code', 'Experience', 'Mortgage'], axis=1)`
`data1`

Out[14]:

	ID	Age	Income	Family	CCAvg	Education	Personal.Loan	CD.Account	Online	CreditCard
0	1	25	49	4	1.6	1	0	0	0	0
1	2	45	34	3	1.5	1	0	0	0	0
2	3	39	11	1	1.0	1	0	0	0	0
3	4	35	100	1	2.7	2	0	0	0	0
4	5	35	45	4	1.0	2	0	0	0	1
...
4995	4996	29	40	1	1.9	3	0	0	1	0
4996	4997	30	15	4	0.4	1	0	0	1	0
4997	4998	63	24	2	0.3	3	0	0	0	0
4998	4999	65	49	3	0.5	2	0	0	1	0
4999	5000	28	83	3	0.8	1	0	0	1	1

5000 rows × 10 columns

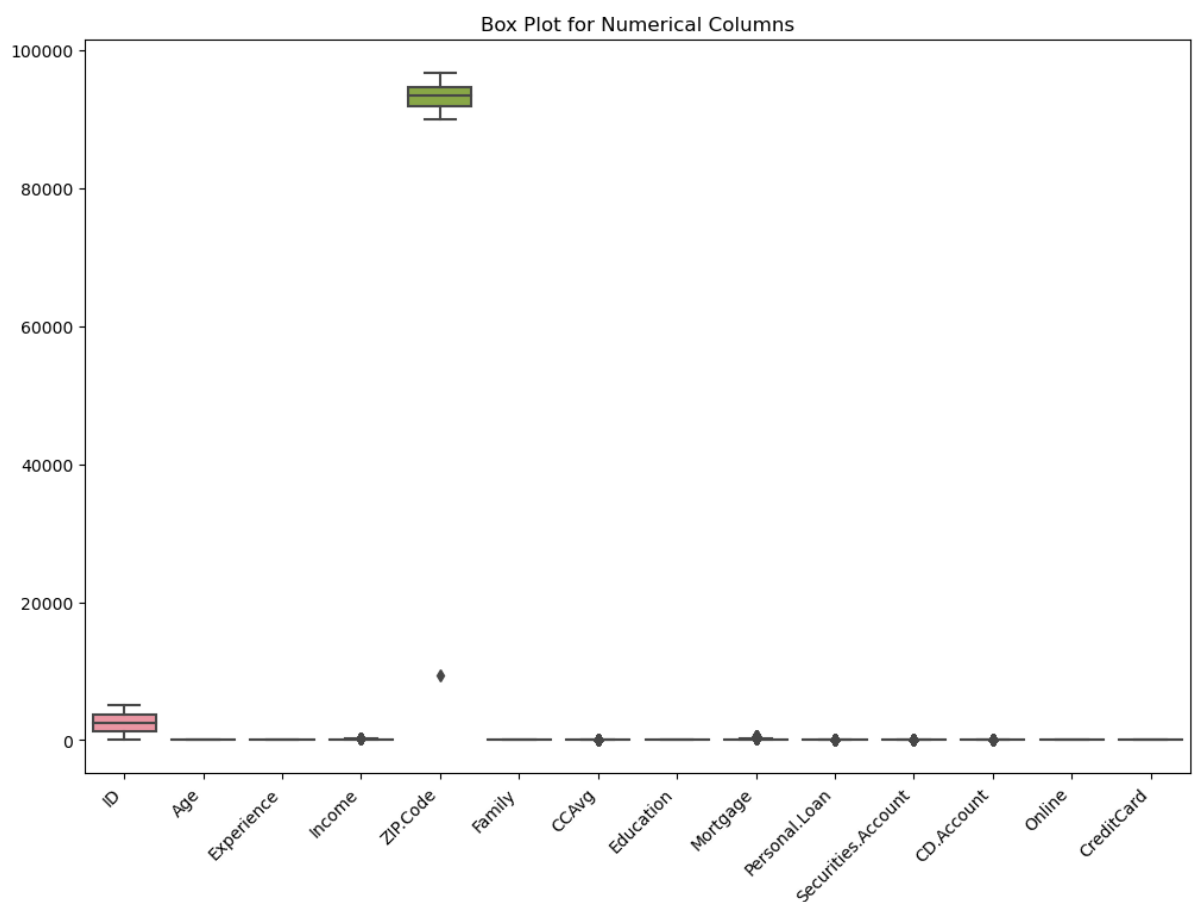
```
In [15]: numerical_columns = data.select_dtypes(include=['number'])

# Create a box plot for all numerical columns

plt.figure(figsize=(12, 8))
sns.boxplot(data=numerical_columns)
plt.title("Box Plot for Numerical Columns")

# Rotate x-axis labels for better visibility
plt.xticks(rotation=45, ha="right")

# Show the plot
plt.show()
```



```
In [16]: ## as we look at Box plots we have outliers in Zip_code column. Thus We Won't use it
## then we will drop it
```

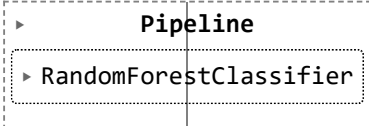
```
In [17]: # Assuming 'ZIP_Code' and 'Personal_Loan' are columns in your DataFrame
x = data.drop(['ZIP.Code', 'Personal.Loan', 'ID'], axis=1) # Drop the specified columns
y = data['Personal.Loan'] # Set the target variable
```

```
In [18]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)
```

```
In [19]: # Create a pipeline
pipeline_rf = Pipeline([
    ('classifier', RandomForestClassifier())
])
```

```
In [20]: # Fit the pipeline
pipeline_rf.fit(x_train, y_train)
```

```
Out[20]:
```



```
  ▸ Pipeline
    ▸ RandomForestClassifier
```

```
In [21]: # Make predictions
y_pred_rf = pipeline_rf.predict(x_test)

# Evaluate the performance
accuracy_rf = accuracy_score(y_test, y_pred_rf)

print("Random Forest Accuracy:", accuracy_rf)
```

Random Forest Accuracy: 0.99

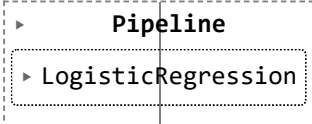
```
In [ ]:
```

```
In [22]: from sklearn.linear_model import LogisticRegression

pipeline_lr = Pipeline([
    ('classifier', LogisticRegression())
])
```

```
In [23]: pipeline_lr.fit(x_train, y_train)
```

```
Out[23]:
```



```
  ▸ Pipeline
    ▸ LogisticRegression
```

```
In [24]: y_pred_lr = pipeline_lr.predict(x_test)

accuracy_lr = accuracy_score(y_test, y_pred_lr)

print("Logistic Regression Accuracy:", accuracy_lr)
```

Logistic Regression Accuracy: 0.952

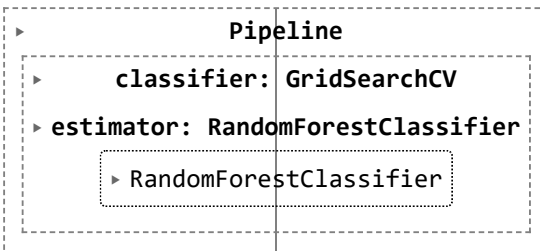
```
In [25]: ## we will use GridSearchCV and will print Best parameters can get Higher Performance
```

```
In [26]: grid_rf = {
    'n_estimators': [10, 50, 100],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
}
```

```
In [27]: # Create a pipeline with GridSearchCV for RandomForestClassifier
pipeline_rf_cv = Pipeline([
    ('classifier', GridSearchCV(RandomForestClassifier(), grid_rf, cv=5))
])
```

```
In [29]: #Fit the pipeline with cross-validation and hyperparameter tuning
pipeline_rf_cv.fit(x_train, y_train)
```

```
Out[29]:
```



```

  Pipeline
  ├── classifier: GridSearchCV
  │   └── estimator: RandomForestClassifier
  │       └── RandomForestClassifier

```

```
In [32]: y_pred_rf_cv = pipeline_rf_cv.predict(x_test)

accuracy_rf_cv = accuracy_score(y_test, y_pred_rf_cv)
print("Random Forest Accuracy (with CV):", accuracy_rf_cv)
```

Random Forest Accuracy (with CV): 0.99

```
In [33]: best_params_rf = pipeline_rf_cv.named_steps['classifier'].best_params_
print("\nBest Hyperparameters for RandomForestClassifier:")
print(best_params_rf)
```

Best Hyperparameters for RandomForestClassifier:
{ 'max_depth': 10, 'min_samples_split': 5, 'n_estimators': 50 }

```
In [34]: # Fit the pipeline on the training data
pipeline_rf_cv.fit(x_train, y_train)

# Make predictions on the training set
y_pred_train = pipeline_rf_cv.predict(x_train)

# Make predictions on the test set
y_pred_test = pipeline_rf_cv.predict(x_test)

# Calculate accuracy for training set
accuracy_train = accuracy_score(y_train, y_pred_train)
print("Training Set Accuracy:", accuracy_train)

# Calculate accuracy for test set
accuracy_test = accuracy_score(y_test, y_pred_test)
print("Test Set Accuracy:", accuracy_test)
```

Training Set Accuracy: 0.99825
Test Set Accuracy: 0.989

```
In [ ]:
```

